

Spark SQL Guide

- Getting Started
- Data Sources
- Performance Tuning
- Distributed SQL Engine
- PySpark Usage Guide for Pandas with Apache Arrow
- Migration Guide
- SQL Reference
 - ANSI Compliance
 - Data Types
 - Datetime Pattern
 - Number Pattern
 - Functions
 - Identifiers
 - Literals
 - Null Semantics
 - SQL Syntax

Data Types

Supported Data Types

Spark SQL and DataFrames support the following data types:

- Numeric types
 - byteType: Represents 1-byte signed integer numbers. The range of numbers is from -128 to 127.
 - shortType: Represents 2-byte signed integer numbers. The range of numbers is from -32768 to 32767.
 - integerType: Represents 4-byte signed integer numbers. The range of numbers is from -2147483648 to 2147483647.
 - longType: Represents 8-byte signed integer numbers. The range of numbers is from -9223372036854775808 to 9223372036854775807.
 - floatType: Represents 4-byte single-precision floating point numbers.
 - doubleType: Represents 8-byte double-precision floating point numbers.
 - decimalType: Represents arbitrary-precision signed decimal numbers. Backed internally by java.math.BigDecimal. A BigDecimal consists of an arbitrary precision integer unscaled value and a 32-bit integer scale.
- String type
 - stringType: Represents character string values.
 - varcharType(length): A variant of stringType which has a length limitation. Data writing will fail if the input string exceeds the length limitation. Note: this type can only be used in table schema, not functions/operators.
 - charType(length): A variant of varcharType(length) which is fixed length. Reading column of type charType(n) always returns string values of length n. Char type column comparison will pad the short one to the longer length.
- Binary type
 - binaryType: Represents byte sequence values.
- Boolean type
 - booleanType: Represents boolean values.
- Datetime type
 - timestampType: Represents values comprising values of fields year, month, day, hour, minute, and second, with the session local time-zone. The timestamp value represents an absolute point in time.
 - dateType: Represents values comprising values of fields year, month and day, without a time-zone.
- Interval types
 - yearMonthIntervalType(startField, endField): Represents a year-month interval which is made up of a contiguous subset of the following fields:
 - MONTH, months within years [0..11],
 - YEAR, years in the range [0..178956970].

Individual interval fields are non-negative, but an interval itself can have a sign, and be negative.

startField is the leftmost field, and endField is the rightmost field of the type. Valid values of startField and endField are 0(MONTH) and 1(YEAR). Supported year-month interval types are:

Year-Month Interval Type	SQL type	An instance of the type
YearMonthIntervalType(YEAR, YEAR) or YearMonthIntervalType(YEAR)	INTERVAL YEAR	INTERVAL '2021' YEAR
YearMonthIntervalType(YEAR, MONTH)	INTERVAL YEAR TO MONTH	INTERVAL '2021-07' YEAR TO MONTH
YearMonthIntervalType(MONTH, MONTH) or YearMonthIntervalType(MONTH)	INTERVAL MONTH	INTERVAL '10' MONTH

- dayTimeIntervalType(startField, endField): Represents a day-time interval which is made up of a contiguous subset of the following fields:
 - SECOND, seconds within minutes and possibly fractions of a second [0..59.999999],
 - MINUTE, minutes within hours [0..59],
 - HOURL, hours within days [0..23],
 - DAY, days in the range [0..106751991].

Individual interval fields are non-negative, but an interval itself can have a sign, and be negative.

startField is the leftmost field, and endField is the rightmost field of the type. Valid values of startField and endField are 0 (DAY), 1 (HOUR), 2 (MINUTE), 3 (SECOND). Supported day-time interval types are:

Day-Time Interval Type	SQL type	An instance of the type
DayTimeIntervalType(DAY, DAY) or DayTimeIntervalType(DAY)	INTERVAL DAY	INTERVAL '100' DAY

Spark SQL Guide

- Getting Started
- Data Sources
- Performance Tuning
- Distributed SQL Engine
- PySpark Usage Guide for Pandas with Apache Arrow
- Migration Guide
- SQL Reference
 - ANSI Compliance
 - Data Types
 - Datetime Pattern
 - Number Pattern
 - Functions
 - Identifiers
 - Literals
 - Null Semantics
 - SQL Syntax

«

Day-Time Interval Type	SQL type	An instance of the type
DayTimeIntervalType(DAY, HOUR)	INTERVAL DAY TO HOUR	INTERVAL '100 10' DAY TO HOUR
DayTimeIntervalType(DAY, MINUTE)	INTERVAL DAY TO MINUTE	INTERVAL '100 10:30' DAY TO MINUTE
DayTimeIntervalType(DAY, SECOND)	INTERVAL DAY TO SECOND	INTERVAL '100 10:30:40.999999' DAY TO SECOND
DayTimeIntervalType(HOUR, HOUR) or DayTimeIntervalType(HOUR)	INTERVAL HOUR	INTERVAL '123' HOUR
DayTimeIntervalType(HOUR, MINUTE)	INTERVAL HOUR TO MINUTE	INTERVAL '123:10' HOUR TO MINUTE
DayTimeIntervalType(HOUR, SECOND)	INTERVAL HOUR TO SECOND	INTERVAL '123:10:59' HOUR TO SECOND
DayTimeIntervalType(MINUTE, MINUTE) or DayTimeIntervalType(MINUTE)	INTERVAL MINUTE	INTERVAL '1000' MINUTE
DayTimeIntervalType(MINUTE, SECOND)	INTERVAL MINUTE TO SECOND	INTERVAL '1000:01.001' MINUTE TO SECOND
DayTimeIntervalType(SECOND, SECOND) or DayTimeIntervalType(SECOND)	INTERVAL SECOND	INTERVAL '1000.000001' SECOND

- Complex types
 - ArrayType(elementType, containsNull): Represents values comprising a sequence of elements with the type of elementType. containsNull is used to indicate if elements in a ArrayType value can have null values.
 - MapType(keyType, valueType, valueContainsNull): Represents values comprising a set of key-value pairs. The data type of keys is described by keyType and the data type of values is described by valueType. For a MapType value, keys are not allowed to have null values. valueContainsNull is used to indicate if values of a MapType value can have null values.
 - StructType(fields): Represents values with the structure described by a sequence of StructFields (fields).
 - StructField(name, dataType, nullable): Represents a field in a StructType. The name of a field is indicated by name. The data type of a field is indicated by dataType. nullable is used to indicate if values of these fields can have null values.

Scala	Java	Python	R	SQL
-----------------------	----------------------	------------------------	-------------------	---------------------

All data types of Spark SQL are located in the package `org.apache.spark.sql.types`. You can access them by doing

```
import org.apache.spark.sql.types._
```

Find full example code at "examples/src/main/scala/org/apache/spark/examples/sql/SparkSQLExample.scala" in the Spark repo.

Data type	Value type in Scala	API to access or create a data type
ByteType	Byte	ByteType
ShortType	Short	ShortType
IntegerType	Int	IntegerType
LongType	Long	LongType
FloatType	Float	FloatType
DoubleType	Double	DoubleType
DecimalType	java.math.BigDecimal	DecimalType
StringType	String	StringType
BinaryType	Array[Byte]	BinaryType
BooleanType	Boolean	BooleanType
TimestampType	java.sql.Timestamp	TimestampType
DateType	java.sql.Date	DateType

Spark SQL Guide

- [Getting Started](#)
- [Data Sources](#)
- [Performance Tuning](#)
- [Distributed SQL Engine](#)
- [PySpark Usage Guide for Pandas with Apache Arrow](#)
- [Migration Guide](#)
- [SQL Reference](#)
 - [ANSI Compliance](#)
 - [Data Types](#)
 - [Datetime Pattern](#)
 - [Number Pattern](#)
 - [Functions](#)
 - [Identifiers](#)
 - [Literals](#)
 - [Null Semantics](#)
 - [SQL Syntax](#)

«

Data type	Value type in Scala	API to access or create a data type
YearMonthIntervalType	java.time.Period	YearMonthIntervalType
DayTimeIntervalType	java.time.Duration	DayTimeIntervalType
ArrayType	scala.collection.Seq	ArrayType(<i>elementType</i> , [<i>containsNull</i>]) Note: The default value of <i>containsNull</i> is true.
MapType	scala.collection.Map	MapType(<i>keyType</i> , <i>valueType</i> , [<i>valueContainsNull</i>]) Note: The default value of <i>valueContainsNull</i> is true.
StructType	org.apache.spark.sql.Row	StructType(<i>fields</i>) Note: <i>fields</i> is a Seq of StructFields. Also, two fields with the same name are not allowed.
StructField	The value type in Scala of the data type of this field(For example, Int for a StructField with the data type IntegerType)	StructField(<i>name</i> , <i>dataType</i> , [<i>nullable</i>]) Note: The default value of <i>nullable</i> is true.

Floating Point Special Values

Spark SQL supports several special floating point values in a case-insensitive manner:

- Inf/+Inf/Infinity/+Infinity: positive infinity
 - `FloatType`: equivalent to `Scala Float.PositiveInfinity`.
 - `DoubleType`: equivalent to `Scala Double.PositiveInfinity`.
- -Inf/-Infinity: negative infinity
 - `FloatType`: equivalent to `Scala Float.NegativeInfinity`.
 - `DoubleType`: equivalent to `Scala Double.NegativeInfinity`.
- NaN: not a number
 - `FloatType`: equivalent to `Scala Float.NaN`.
 - `DoubleType`: equivalent to `Scala Double.NaN`.

Positive/Negative Infinity Semantics

There is special handling for positive and negative infinity. They have the following semantics:

- Positive infinity multiplied by any positive value returns positive infinity.
- Negative infinity multiplied by any positive value returns negative infinity.
- Positive infinity multiplied by any negative value returns negative infinity.
- Negative infinity multiplied by any negative value returns positive infinity.
- Positive/negative infinity multiplied by 0 returns NaN.
- Positive/negative infinity is equal to itself.
- In aggregations, all positive infinity values are grouped together. Similarly, all negative infinity values are grouped together.
- Positive infinity and negative infinity are treated as normal values in join keys.
- Positive infinity sorts lower than NaN and higher than any other values.
- Negative infinity sorts lower than any other values.

NaN Semantics

There is special handling for not-a-number (NaN) when dealing with `float` or `double` types that do not exactly match standard floating point semantics. Specifically:

- `NaN = NaN` returns true.
- In aggregations, all NaN values are grouped together.
- NaN is treated as a normal value in join keys.
- NaN values go last when in ascending order, larger than any other numeric value.

Examples

Spark SQL Guide

- [Getting Started](#)
- [Data Sources](#)
- [Performance Tuning](#)
- [Distributed SQL Engine](#)
- [PySpark Usage Guide for Pandas with Apache Arrow](#)
- [Migration Guide](#)
- [SQL Reference](#)
 - [ANSI Compliance](#)
 - [Data Types](#)
 - [Datetime Pattern](#)
 - [Number Pattern](#)
 - [Functions](#)
 - [Identifiers](#)
 - [Literals](#)
 - [Null Semantics](#)
 - [SQL Syntax](#)

«

```
SELECT double('infinity') AS col;
+-----+
|      col|
+-----+
|Infinity|
+-----+

SELECT float('-inf') AS col;
+-----+
|      col|
+-----+
| -Infinity|
+-----+

SELECT float('NaN') AS col;
+----+
| col|
+----+
|NaN|
+----+

SELECT double('infinity') * 0 AS col;
+----+
| col|
+----+
|NaN|
+----+

SELECT double('-infinity') * (-1234567) AS col;
+-----+
|      col|
+-----+
|Infinity|
+-----+

SELECT double('infinity') < double('NaN') AS col;
+----+
| col|
+----+
|true|
+----+

SELECT double('NaN') = double('NaN') AS col;
+----+
| col|
+----+
|true|
+----+

SELECT double('inf') = double('infinity') AS col;
+----+
| col|
+----+
|true|
+----+

CREATE TABLE test (c1 int, c2 double);
INSERT INTO test VALUES (1, double('infinity'));
INSERT INTO test VALUES (2, double('infinity'));
INSERT INTO test VALUES (3, double('inf'));
INSERT INTO test VALUES (4, double('-inf'));
INSERT INTO test VALUES (5, double('NaN'));
INSERT INTO test VALUES (6, double('NaN'));
INSERT INTO test VALUES (7, double('-infinity'));
SELECT COUNT(*), c2 FROM test GROUP BY c2;
+-----+-----+
| count(1)|      c2|
+-----+-----+
|         2|      NaN|
|         2| -Infinity|
|         3|  Infinity|
+-----+-----+
```